

Industrial Organization II Referee Report

Sebastian Rodriguez

6 May 2026

Cover letter to editor

This is a referee report of the working paper, *Where to Build Affordable Housing? Evaluating the Tradeoffs of Location* by authors Cody Cook, Pearl Z. Li, and Ariel J. Binder. The paper looks at the tradeoffs of where to locate affordable housing, examining specifically the units built out of the Low-Income Housing Tax Credit program (LIHTC). The main insight of the paper is that choosing a location also determines which tenants receive that assistance. Despite the LIHTC having uniform eligibility rules within cities, new developments built in areas with better opportunities (schools, jobs, etc) end up serving households that are whiter, more educated, have fewer children and have a higher income within the eligibility pool as compared to new developments built in lower-opportunity areas. The paper builds a structural choice model to quantify the tradeoffs.

I find that the paper's argument is convincing and believe it is a strong fit for the journal. My main concerns are in the credibility of the magnitudes of the welfare results. I believe that these concerns can be addressed through additional robustness analysis and discussion. I recommend a revise and resubmit.

Summary of paper

In their paper, the authors combine IRS tax records, Census migration histories, ACS data and HUD administrative data on LIHTC properties to assemble data on LIHTC households across the 50 most populated metropolitan statistical areas from 2010-2018. The paper finds that there are large differences across neighborhoods in the characteristics of LIHTC tenants. LIHTC households in the top quartile of neighborhood opportunity are twice as likely to have been college educated, and are three times more likely to be white as compared to the bottom quartile. To interpret these patterns, the authors constructed a static residential choice model where households decide whether to apply to an affordable housing unit and then developers allocate the units using a weighted lottery. Market-rate preferences for housing and neighborhood characteristics are estimated using a Bayer-Ferreira-McMillan approach. Household preferences specific to affordable housing and developer lottery weights are identified by GMM specifically using the assumption that developers affect move-ins and not move-outs.

The authors estimate that building a new unit in a high-opportunity neighborhood raises the aggregate household surplus by 151 dollars a month per unit but shifts the distribution of beneficiaries away from Black and Hispanic households towards white households. This corresponds to a loss

of 273 dollars per month in expected surplus. Additionally, they report that 76% of the decline in the Black/Hispanic share comes from the increase in applications from white applicants.

Assessment of contribution

The main contribution of this paper is both empirical and conceptual. This paper shows that where you build affordable housing determines who applies for it. When developers choose to build in higher opportunity areas, moderate-need white households' applications crowd out higher-need minority households. The authors build a structural model to quantify this, finding that these higher opportunity developments raise aggregate welfare but actually hurt Black and Hispanic households in expected surplus terms. This is important since affordable housing policy has shifted toward higher-opportunity neighborhoods. This policy change assumed that this would grant the neediest households better access to schools and jobs. This contributes to ongoing policy debates about where to place LIHTC developments. If policy makers want the opportunity benefits to reach highest need minorities, then these programs need better targeting of these disadvantaged households.

Major comments

1. My first concern related to the data used. A core feature of the argument is that moderate need applicants crowd out high-need applicants. Unfortunately, the authors do not have any data on who applies for the LIHTC units. The authors instead infer application behavior from the observed allocations and moves. Their model must discern household preferences for different locations from the allocation mechanism used by the developers since both determine who ultimately gets the unit. They use move-out behavior to separately identify household preferences. While this is a clever move, without the direct data on applications, it is unclear to what extent the model can identify the two forces rather than attributing patterns in allocation to one force versus the other. Using move out preferences to estimate move in preferences assumes that households move out for reasons unrelated to what developer they are with. We can imagine that move out behavior is partly caused by the developers screening move ins. I think further analyses on the model's sensitivity to alternative assumptions about application decisions will provide additional evidence supporting their argument.
2. Another concern was the simplification of how developers choose who are allocated units. In the model, the LIHTC units are allocated based on a weighted lottery where the developers assign weights based on household characteristics. They argue that this reduced form approximation captures any allocation mechanism where applicants face equal offer probabilities. They justify this simplification by showing that for-profit and non-profit developers have sim-

ilar tenant composition. While this does allow for a computable way to model the rationing, the authors acknowledge that the actual allocation processes are much more heterogeneous in practice. This includes wait-lists, first come first serve rules, referrals and varying degrees of screening by developers. The key results of the paper depend critically on how applicants are selected, so if the true processes are significantly different than those in the weighted lottery, then the estimated model may misattribute some allocation patterns to household preferences rather than institutional features. I think exploring alternative allocation mechanisms could demonstrate that the results are not being driven by the chosen form of the allocation process.

3. Another concern I will point to is that the authors estimate household preferences for housing and neighborhood characteristics using choices observed in the normal private sector. They then apply these estimated preferences to households applying to the LIHTC units. This approach relies on the assumption that households in market-rate and LIHTC units share the same observed preferences. The assumption that these preferences generalize to the population of interest may be quite strong. These characteristics may be correlated with preferences for neighborhood characteristics. The authors defend this by comparing LIHTC households with households who were eligible for LIHTC but didn't end up in a unit. After controlling for observables, they find that the difference between the two groups is very small. However, we can imagine that the households that are picked differ systematically from the households who are not picked along unobserved dimensions such as search frictions, credit constraints and risk tolerance. The models counterfactual results directly depend on these estimated preferences so if they are mismeasured, the welfare comparison results may be biased. I think adding additional validation for this assumption is necessary.

Minor comments and suggestions

1. The structural model is estimated only with data from Chicago. Chicago seems like a specific case where there is high segregation and a particular housing stock. I think the paper should specifically discuss how the Chicago specific estimation influences the results.
2. I think many of the tables included in the appendices should be featured in the main description section. Specifically, table D.4.
3. The neighborhood opportunity index features five components of equal weights. It would be interesting to see what components are driving the variation.

Recommendation

My recommendation is that the authors revise and resubmit. The paper's contribution is important and has a novel empirical strategy. The major comments are correctable as most require additional sensitivity analyses and additional discussion. If these concerns are fixed, I believe this paper should be accepted.